Fleet Policy Q&A with Citations - Architecture

Single HTTP endpoint in front of an Azure AI Foundry agent (GPT-5 + File Search). The agent retrieves and reasons; the service enforces the output contract.

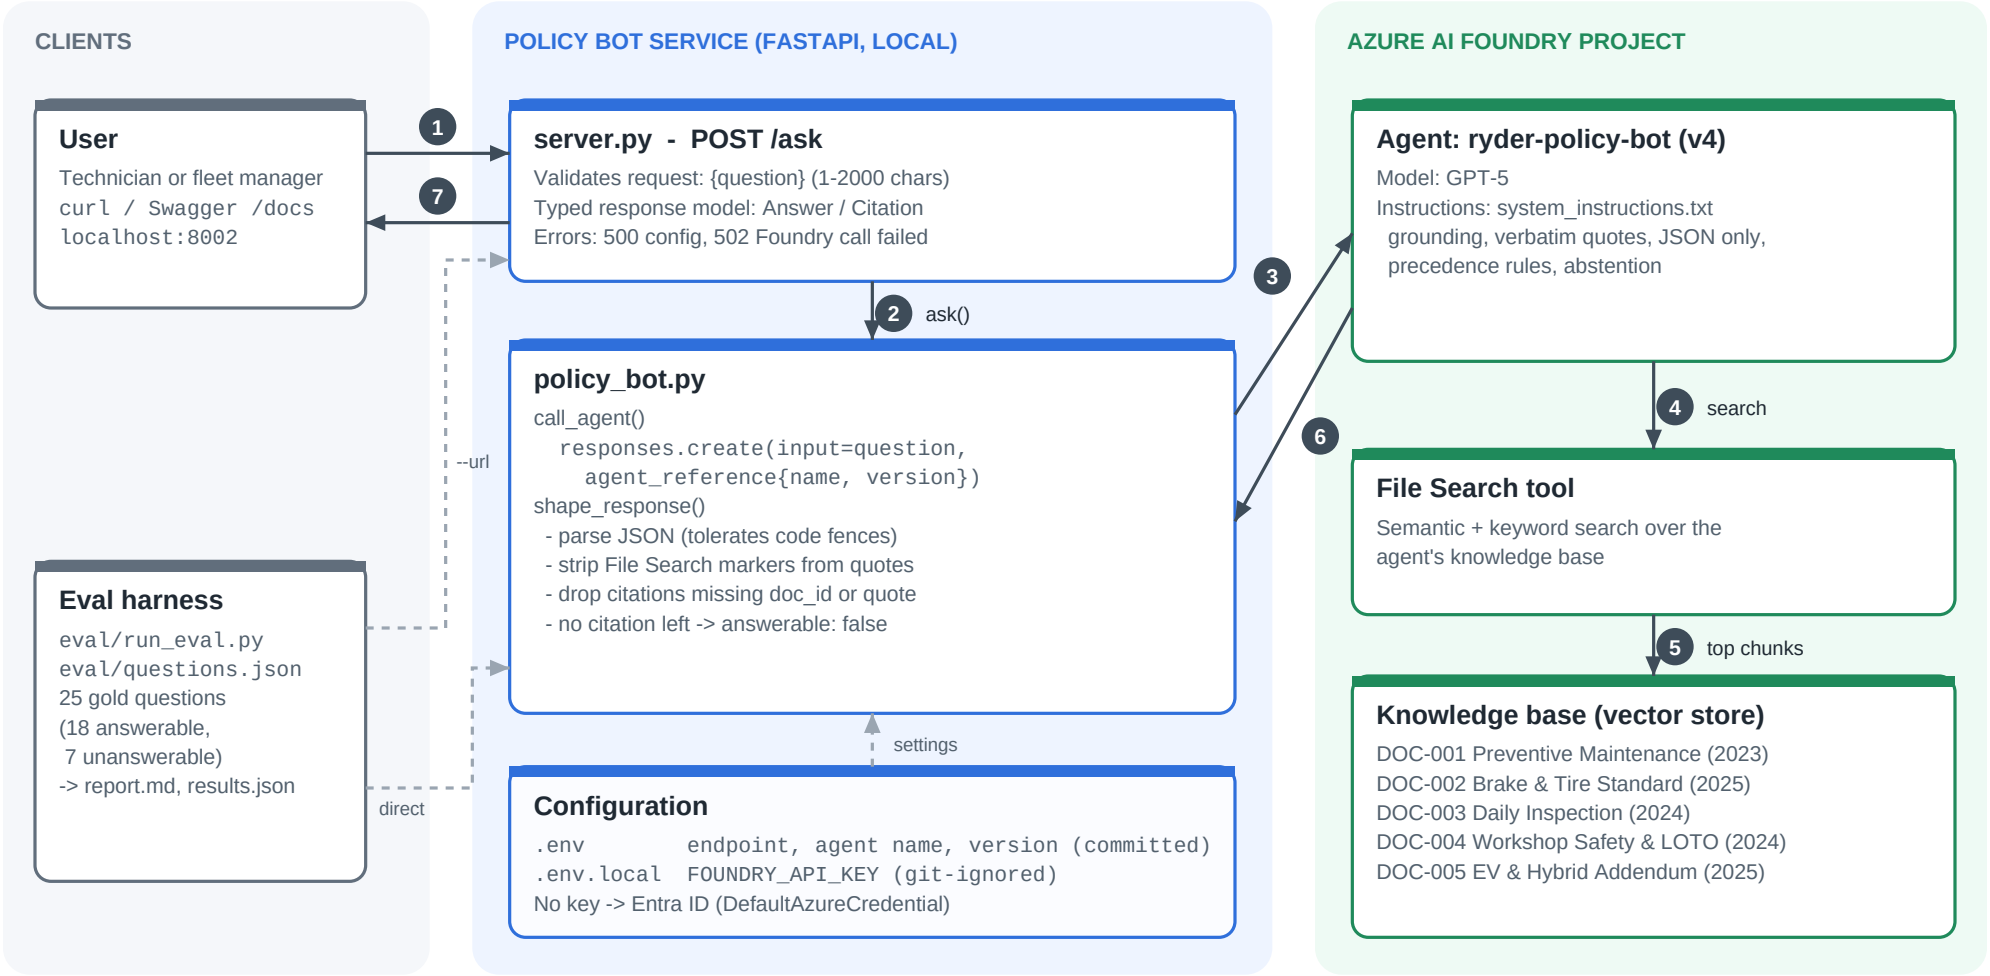

Request flow

1 Client sends a question
2 Server hands it to policy_bot.ask()
3 Responses API call to the agent (api-key header)
4 Agent runs File Search
5 Relevant policy passages come back
6 Agent replies with JSON text
7 Contract-checked answer returned

Output contract (every response)

```
{ "question": str,
  "answer": str,
  "citations": [{"doc_id", "section", "quote"}],
  "answerable": bool }
```

answerable: true always has >= 1 citation.
Not in the corpus -> answerable: false, no citations.
Conflicts: later effective date / EV addendum prevail.

Measured (agent v4, 25 questions)

Overall accuracy	**100%**
Section recall	**100%**
Citation precision	**96%**
Abstention recall	**100%**
Latency (median / p90)	**7.8 s / 14.1 s**
Cost per request	**~ $0.016**